

On the relationship between logic and databases

The case of relational database queries

Maurizio Lenzerini

Sapienza Università di Roma

A.Y. 2021/22

Example of interpretation

Example (of first-order language and interpretation)

Symbols Constants: 1, 2, 3, 4, 5
Predicate symbols: *has-mother*/2, *friends*/2

Domain $\Delta = \{1, 2, 3, 4, 5\}$

Interpretation $\mathcal{I}(1) = 1, \mathcal{I}(2) = 2, \mathcal{I}(3) = 3, \mathcal{I}(4) = 4, \mathcal{I}(5) = 5$

$\mathcal{I}(\textit{has-mother}) = \left\{ \begin{array}{l} \langle 1, 2 \rangle, \langle 2, 3 \rangle \\ \langle 3, 4 \rangle, \langle 4, 5 \rangle \end{array} \right\}$

$\mathcal{I}(\textit{friends}) = \left\{ \begin{array}{l} \langle 1, 2 \rangle, \langle 2, 1 \rangle, \langle 3, 4 \rangle, \\ \langle 4, 3 \rangle, \langle 4, 2 \rangle, \langle 2, 4 \rangle, \\ \langle 4, 1 \rangle, \langle 1, 4 \rangle, \langle 4, 4 \rangle \end{array} \right\}$

It is immediate to see the above first-order interpretation as a relational database, where the domain determines the possible values in the database, each predicate corresponds to a relation, and the extension of each predicate determines the tuples in the corresponding relation.

Database as a first-order interpretation

Let D be a relational database, and let Δ , called the **active domain** of D , be the set of values stored in D . Since D is finite, Δ is also finite. We often use $adom(D)$ to denote Δ .

Let us define a first-order language L_D with equality (and other “interpreted” predicates) as follows:

- the set of constant symbols is simply Δ ,
- the set of function symbols is empty,
- the set of predicate symbols includes
 - one symbol P/n for each relation P with n columns (where argument i corresponds to the i -th attribute),
 - the equality predicate, plus other predicates: $\neq, <, \leq, \dots$

It is easy to see that D is a **finite interpretation for L_D** such that:

- the **domain** of such interpretation is Δ , and therefore is finite
- every constant is mapped to itself, and therefore different constants are interpreted as different domain elements (unique name assumption)
- the **interpretation function** is given by the extension of the various relations in D , where each of them is finite

First-order language and query language

Since the database D plays the role of interpretation for L_D , the formulas of the language L_D can now be evaluated with respect to such an interpretation.

This is the basic idea of “**logic as a query language**”: an open formula in L_D with free variables x_1, \dots, x_k will correspond to a query that, when evaluated with respect to D , will return the k -tuples of constants in D that, when assigned to the variables x_1, \dots, x_k , make the formula true.

In other words, the formula defines a relation over the database D , which is the result of the query.

The idea was proposed by E. Codd, the inventor of the relational model. In contrast to the **relational algebra**, also proposed by Codd, the query language so defined was called **relational calculus** (with the semantics based on the active domain).

Relational calculus expressions

We will now see, given D , how we will write queries as relational calculus expressions in L_D

A **relational calculus expression** in L_D is an expression of the form

$$\{(x_1, \dots, x_k) : \phi(x_1, \dots, x_k)\}$$

where $\phi(x_1, \dots, x_k)$ is a first-order formula of L_D with x_1, \dots, x_k as its free variables.

When applied to a relational database D , this relational calculus expression returns the k -ary relation that consists of all k -tuples (a_1, \dots, a_k) of constants in D that make the formula true on D .

Relational calculus expressions

Consider the interpretation/database shown at page 2.

Example

The relational calculus expression

$$\{(x, y) : \exists z(\textit{has-mother}(x, z) \wedge \textit{has-mother}(z, y))\}$$

returns the set of all pairs (a, b) such that b is the mother of the mother of a .

Example

The relational calculus expression

$$\{(x) : \exists y(\textit{has-mother}(x, y) \wedge \forall z(\textit{friends}(x, z) \rightarrow \textit{friends}(y, z)))\}$$

returns the set of all objects a all of whose friends are friends of her/his mother.

Relationship between relational algebra and calculus

If Q is a query and D a database, let us denote by $Q(D)$ the result of evaluating Q wrt D .

Theorem (Codd's theorem)

The relational algebra and the relational calculus are “equivalent”, i.e.,

- *For every relational algebra expression E there is a relational calculus query F such that for every database D , $E(D) = F(D)$.*
- *For every relational calculus query F there is a relational algebra expression E such that for every database D , $E(D) = F(D)$.*

Proof of Codd's theorem – from algebra to calculus: By induction on the relational algebra expression.

We have a relational database whose schema is `Movie(title, director, actor)` and `Schedule(theater, mtitle)`, where both attributes, title and mtitle, refer to the title of a movie.

Consider the following queries:

- Which theaters show some movies directed by Tarantino?
- Which theaters do not show any movies directed by Tarantino?
- Which theaters show only movies directed by Tarantino?
- Which theaters show all movies directed by Tarantino?

Express each of the queries above in the three query languages of

- Relational Calculus
- Relational Algebra
- SQL

Exercise

We have a relational database whose schema is `Movie(title, director, actor)` and `Schedule(theater, mtitle)`, where both attributes, title and mtitle, refer to the title of a movie.

Consider the following queries:

- Which theaters show some movies directed by Tarantino?
 $\{(x) \mid \exists y \exists z \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, z)\}$
- Which theaters do not show any movies directed by Tarantino?
- Which theaters show only movies directed by Tarantino?
- Which theaters show all movies directed by Tarantino?

We have a relational database whose schema is **Movie**(title, director, actor) and **Schedule**(theater, mtitle), where both attributes, title and mtitle, refer to the title of a movie.

Consider the following queries:

- Which theaters show some movies directed by Tarantino?
 $\{(x) \mid \exists y \exists z \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, z)\}$
- Which theaters do not show any movies directed by Tarantino?
 $\{(x) \mid \neg \exists y \exists w \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, w)\}$
- Which theaters show only movies directed by Tarantino?
- Which theaters show all movies directed by Tarantino?

We have a relational database whose schema is **Movie**(title, director, actor) and **Schedule**(theater, mtitle), where both attributes, title and mtitle, refer to the title of a movie.

Consider the following queries:

- Which theaters show some movies directed by Tarantino?
 $\{(x) \mid \exists y \exists z \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, z)\}$
- Which theaters do not show any movies directed by Tarantino?
 $\{(x) \mid \neg \exists y \exists w \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, w)\}$
- Which theaters show only movies directed by Tarantino?
 $\{(x) \mid \forall y (\text{Schedule}(x, y) \rightarrow \exists w \text{ Movie}(y, \text{Tarantino}, w))\}$
- Which theaters show all movies directed by Tarantino?

Exercise

We have a relational database whose schema is **Movie**(title, director, actor) and **Schedule**(theater, mtitle), where both attributes, title and mtitle, refer to the title of a movie.

Consider the following queries:

- Which theaters show some movies directed by Tarantino?
 $\{(x) \mid \exists y \exists z \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, z)\}$
- Which theaters do not show any movies directed by Tarantino?
 $\{(x) \mid \neg \exists y \exists w \text{ Schedule}(x, y) \wedge \text{Movie}(y, \text{Tarantino}, w)\}$
- Which theaters show only movies directed by Tarantino?
 $\{(x) \mid \forall y (\text{Schedule}(x, y) \rightarrow \exists w \text{ Movie}(y, \text{Tarantino}, w))\}$
- Which theaters show all movies directed by Tarantino?
 $\{(x) \mid \forall y ((\exists w \text{ Movie}(y, \text{Tarantino}, w)) \rightarrow \text{Schedule}(x, y))\}$

First-order logic and integrity constraints

We recall that an integrity constraints is a condition that the database has to satisfy in order to be coherent with the domain it represents.

It is interesting to observe that logic (in particular, first-order logic) is an ideal language for expressing integrity constraints over a database schema. For examples:

- **Key constraint:** to specify that the attribute A is a key of the relation $R(A, B, C)$, the following formula can be used:

$$\forall x \forall y_1 \forall y_2 \forall z_1 \forall z_2 (R(x, y_1, z_1) \wedge R(x, y_2, z_2) \rightarrow (y_1 = y_2 \wedge y_2 = z_2))$$

- **Foreign key constraint:** if A is a key of $R(A, B, C)$, then to specify that the attribute E in $Q(D, E)$ is a foreign key of R , the following formula can be used:

$$\forall x \forall y (Q(x, y) \rightarrow \exists w \exists z R(y, w, z))$$

More generally, logic allows expressing many sophisticated conditions on the database.

First-order logic and integrity constraints

We remind the reader that a schema S is constituted by the alphabet of the database (fixing the schema of each of the relations) and a set of integrity constraints.

If we express integrity constraints in first-order logic, the notion of database is formalized as follows.

Definition

If $S = \langle \Sigma, \Gamma \rangle$ is a relational schema with alphabet Σ and set Γ of integrity constraints, a database D assigning an extension to all the relations in Σ is said to be **legal** for S if $D \models \Gamma$, i.e., the interpretation D is a model of all the formulas expressing the integrity constraints in Γ .

A database that is legal for S is also called an S -database. In summary, **an S -database is a model of the schema S , where S is seen as a logical theory.**

The complexity of the query evaluation problem

Vardi's taxonomy (M.Y Vardi, "The Complexity of Relational Query Languages", 1982):

Definition

Let L be a database query language.

- The **combined complexity** of L is the complexity of the following decision problem: given as input an L -sentence ϕ and a database D , is ϕ true on D ? (in symbols, $D \models \phi$?)
- The **data complexity** of L is the family of the following decision problems P_ϕ , where ϕ is an L -sentence: given as input a database D , $D \models \phi$?
- The **query complexity** of L is the family of the following decision problems P_D , where D is a database: given as input an L -sentence ϕ , $D \models \phi$?

Complexity of query evaluation in the relational calculus

Definition

- The **combined complexity** of the relational calculus is PSPACE-complete
- The **query complexity** of the relational calculus is PSPACE-complete
- The **data complexity** of the relational calculus is in LOGSPACE (and therefore in PTIME)

Question: Are there interesting sublanguages of relational calculus for which the Query Evaluation Problem are “easier” than the full relational calculus?

Definition

A **conjunctive query** is a query expressible by a relational calculus formula in prenex normal form (i.e., all quantifiers appear at the beginning of the formula) built from atomic formulas $R(y_1, \dots, y_n)$, and \wedge and \exists only:

$$\{(x_1, \dots, x_k) \mid \exists z_1 \dots \exists z_m \chi(x_1, \dots, x_k, z_1, \dots, z_m)\}$$

where $\chi(x_1, \dots, x_k, z_1, \dots, z_m)$ is a conjunction of atomic formulas, called the *body* of the query.

These queries corresponds to relational algebra expressions of the form $PROJ_X(SEL_\gamma(R_1 \times \dots \times R_n))$, where γ is a conjunction of equality atoms.

They also correspond to the basic form of SQL, namely “SELECT FROM WHERE”.

Some notation

- Recall that for a database D , the active domain of D , called $adom(D)$, denotes the values appearing in the relations of D .
- We have observed that a database can be seen as a first-order interpretation. Since the projection of an interpretation function on predicate symbols can be seen as a set of facts (i.e., ground atomic formulas), we can regard a database simply as a set of facts:

$$\{P(a_1, \dots, a_n) \mid (a_1, \dots, a_n) \text{ is a tuple in the relation } P \text{ of } D\}$$

- We often write a conjunction query $\{(x_1, \dots, x_n) \mid \exists z_1 \dots \exists z_k \gamma_1 \wedge \gamma_2 \wedge \dots \wedge \gamma_m\}$ of a conjunction query in the simplified form:

$$\{(x_1, \dots, x_n) \mid \gamma_1, \gamma_2, \dots, \gamma_m\}$$

The notion of homomorphism

Definition (Homomorphism)

Let D and F be two databases over the same relational schema S . A **homomorphism** $h : D \rightarrow F$ is a function from $\text{atom}(D)$ to $\text{atom}(F)$ such that for every relational symbol P of S and every (a_1, \dots, a_m) , we have that if $(a_1, \dots, a_m) \in P^D$, then $(h(a_1), \dots, h(a_m)) \in P^F$.

In what follows, we often write $h((a_1, \dots, a_m))$ to denote $(h(a_1), \dots, h(a_m))$.

Homomorphism: example

- $D_1 = \{P_1(a, b), P_2(b, c, d), P_2(c, a, b)\}$
- $D_2 = \{P_1(n, m), P_1(n, n), P_2(m, m, r), P_2(m, n, m)\}$
- h_1 such that $h_1(a) = n, h_1(b) = m, h_1(c) = m, h_1(d) = r$ is a homomorphism from D_1 to D_2
- no homomorphism from D_2 to D_1 exists.

The homomorphism problem

Definition (Homomorphism problem)

Given two databases D and F , is there a homomorphism $h : D \rightarrow F$?

The homomorphism problem is a fundamental algorithmic problem:

- Satisfiability can be viewed as a special case of it.
- k-Colorability can be viewed as a special case of it.
- Many Artificial Intelligence problems, such as planning, can be viewed as a special case of it.
- Every constraint satisfaction problem can be viewed as a special case of the Homomorphism Problem (Feder and Vardi – 1993).

Duality between databases and conjunctive queries

A Boolean conjunctive query can be seen as a database.

Definition (Canonical database of a conjunctive query)

Given a conjunctive query Q , the **canonical database** of Q is the database D^Q with the variables of Q as active domain and the conjuncts of Q as database facts.

For example, the canonical database of

$\{() \mid E(x, y) \wedge E(y, z) \wedge E(z, w)\}$ is constituted by the facts $\{E(x, y), E(y, z), E(z, w)\}$.

Duality between databases and conjunctive queries

A database can be seen as a Boolean conjunctive query.

Definition (Canonical conjunctive query of a database)

Given a database D , the **canonical conjunctive query** Q^D is the conjunctive query with (a renaming of) the elements in $adom(D)$ as variables and the facts of D as conjuncts.

For example, the canonical conjunctive query of the database $\{E(a, b) \wedge E(b, c) \wedge E(c, a)\}$ is $\{() \mid \exists x \exists y \exists z E(x, y), E(y, z), E(z, x)\}$.

Complexity of conjunction query evaluation

Theorem

Checking whether a Boolean conjunctive query

$$\{() \mid \exists z_1 \dots \exists z_m \chi(z_1, \dots, z_m)\}$$

is true with respect to a database is NP-complete in combined complexity.

An NP algorithm for the problem of checking whether a Boolean conjunctive query Q is true wrt D is as follows:

- 1 guess a function f from $\text{adom}(D^Q)$ to $\text{adom}(D)$,
- 2 check in polynomial time if f is a homomorphism from D^Q to D .

Indeed, a function from $\text{adom}(D^Q)$ to $\text{adom}(D)$, corresponds to an assignment to all quantified variables of Q with values in the database D , and then checking whether such a function is a homomorphism corresponds to checking whether the assignment makes the query Q true wrt D .

Complexity of conjunction query evaluation

Theorem

Checking whether a Boolean conjunctive query

$$\{() \mid \exists z_1 \dots \exists z_m \chi(z_1, \dots, z_m)\}$$

is true with respect to a database is NP-complete in combined complexity.

The proof of NP-hardness is by reduction to 3-colorability:

- the complete undirected graph with 3 nodes (such nodes are called (b)lue, (g)reen and (r)ed respectively), can be represented as the database
$$K_3 = \{E(r, b), E(b, r), E(b, g), E(g, b), E(r, g), E(g, r)\}$$
- if G is a graph, then the query Q_G associated to G is the Boolean conjunctive query that has as existential quantified variables all the nodes of G , and as body the conjunction of atoms $E(x, y)$ for each edge of G from x to y
- Q_G is true with respect to K_3 if and only if G is 3-colorable.

Query containment and the containment problem

Definition

Given two queries Q_1 and Q_2 , we say that Q_1 is **contained** in Q_2 , written as $Q_1 \subseteq Q_2$, if for every database D we have that $Q(D) \subseteq Q_2(D)$.

Definition

The **containment problem** is the following decision problem: given two queries Q_1 and Q_2 , decide whether $Q_1 \subseteq Q_2$.

Note that the containment problem is a logical implication problem in *finite models*.

Indeed, if Q_1 and Q_2 have arity n , then $Q_1 \subseteq Q_2$ if and only if

$$\models_{fin} \forall x_1 \forall x_2 \dots \forall x_n \quad Q_1(x_1, \dots, x_n) \rightarrow Q_2(x_1, \dots, x_n)$$

where $\models_{fin} \alpha$ means that α is true in every finite interpretation.

Complexity of conjunctive query containment

Theorem

The containment problem for conjunctive queries, i.e., checking whether $Q_1 \subseteq Q_2$, where Q_1 and Q_2 are conjunctive queries, is an NP-complete problem.

We leave the proof as an exercise. Hint: prove that $Q_1 \subseteq Q_2$ if and only if there is a homomorphism from Q_1^D to Q_2^D .

Unions of conjunctive queries

Definition

A union of conjunctive queries is a query of the form:

$$\{(x_1, \dots, x_n) \mid \gamma_1 \vee \dots \vee \gamma_m\}$$

where each γ_i is the body of a conjunction query with free variables x_1, \dots, x_n .

Exercise

Prove the following:

- The complexity of evaluating unions of conjunctive queries is the same as the complexity of evaluating conjunctive queries.
- The complexity of query containment for unions of conjunctive queries is the same as the complexity of query containment for conjunctive queries.